

What have we learned?

- `fopen` returns a **FILE pointer**: (null pointer if fails)

```
fp = fopen("in.dat", "r");
```

- The `fclose` function closes a file that is no longer in use:

```
int fclose(FILE *fp);
```

Text

00110011	00110010	00110111	00110110	00110111
'3'	'2'	'7'	'6'	'7'

Binary

01111111	11111111
----------	----------

Block I/O

- `fread` and `fwrite` allow a program to **read** and **write** large blocks of data in a single step
- `fread` and `fwrite` are used primarily with **binary streams**, although—**with care**—it's possible to use them with **text** streams

fwrite function

- `fwrite` copies an **array** from memory to a **stream**

array **size of each element (in bytes)**

```
fwrite(a, sizeof(a[0]),  
      sizeof(a) / sizeof(a[0]), fp);
```

number of elements to be copied **FILE pointer**

- returns the number of elements **actually written**

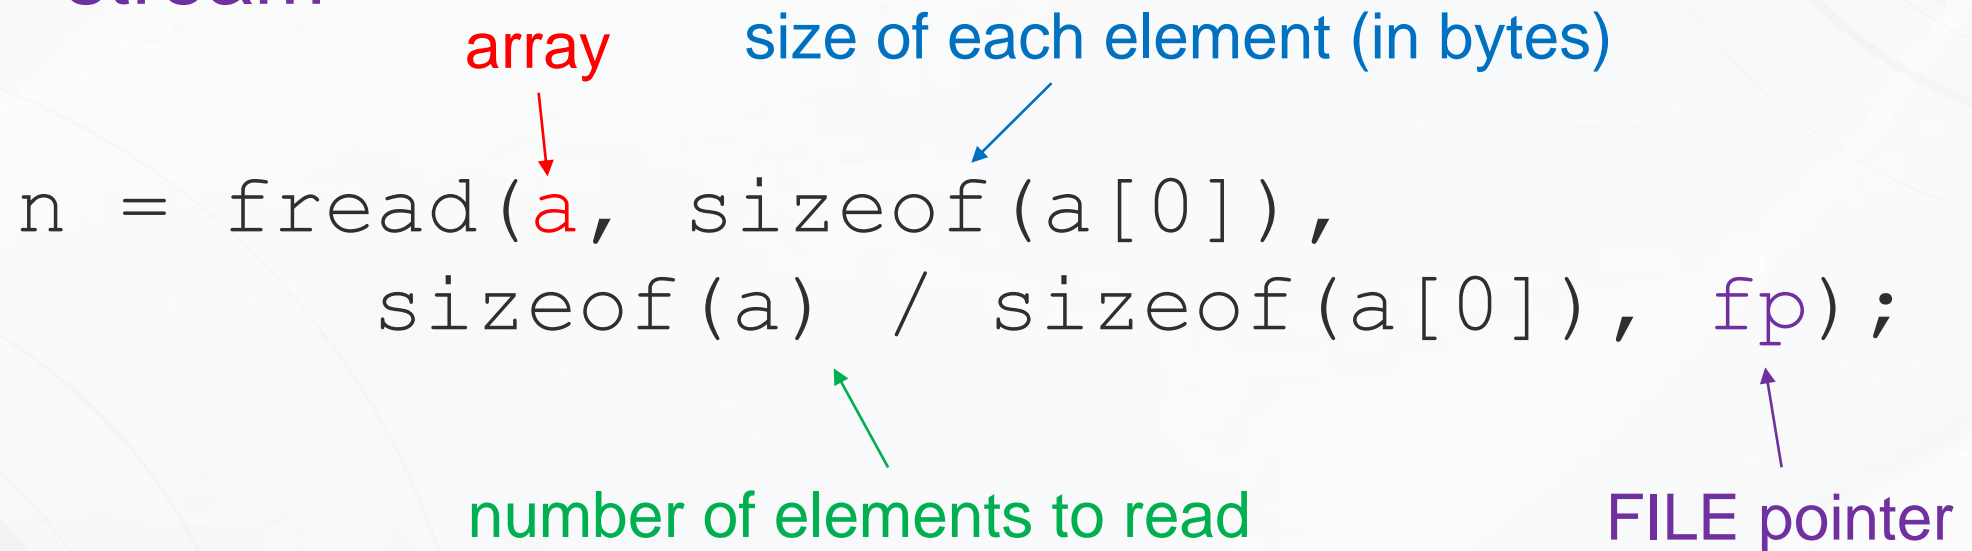
fread function

- `fread` will read the elements of an array from a stream

`n = fread(a, sizeof(a[0]), sizeof(a) / sizeof(a[0]), fp);`

array size of each element (in bytes)

number of elements to read FILE pointer

The diagram illustrates the parameters of the fread function. The first parameter 'a' is red, with a red arrow pointing to it from the label 'array'. The second parameter 'sizeof(a[0])' is blue, with a blue arrow pointing to it from the label 'size of each element (in bytes)'. The third parameter 'sizeof(a) / sizeof(a[0])' is green, with a green arrow pointing to it from the label 'number of elements to read'. The fourth parameter 'fp' is purple, with a purple arrow pointing to it from the label 'FILE pointer'.

- returns the number of elements **actually read**

Block I/O

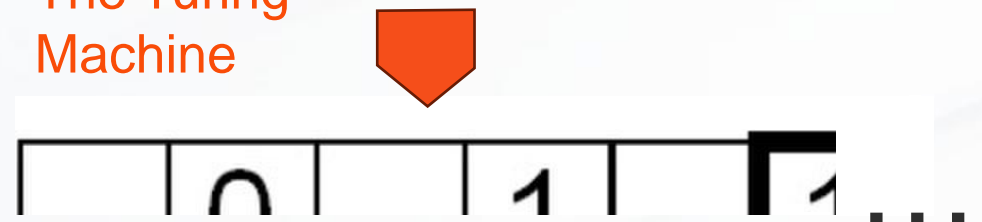
- `fwrite` is convenient to **save data** before terminating a program.
- Later, the program (or another program) can use `fread` to read the **data** back into memory.
- The data doesn't have to be in array form.
- A call of `fwrite` that writes **a variable**/structure/union... `s` to a file:

```
fwrite(&s, sizeof(s), 1, fp);
```

File positioning

- Every stream has an associated *file position*.
- When a file is opened, the file position is set at the *beginning of the file*.
 - In “append” mode, the initial file position may be at the beginning or end, depending on the implementation.
- When a read or write operation is performed, the file *position advances automatically*, providing *sequential access* to data.

The Turing
Machine



File positioning

- `fseek` function changes the **file position** associated with the first argument (a **FILE pointer**).
- The third argument is one of three macros:

<code>SEEK_SET</code>	Beginning of file
<code>SEEK_CUR</code>	Current file position
<code>SEEK_END</code>	End of file
- The second argument, which has type `long int`, is a (possibly negative) byte count, to **offset** from any of the three positions.

File positioning examples

- Using `fseek` to move to the **beginning of a file**:

```
fseek(fp, 0L, SEEK_SET);
```

- Using `fseek` to move to the **end of a file**:

```
fseek(fp, 0L, SEEK_END);
```

- Using `fseek` to **move back 10** bytes:

```
fseek(fp, -10L, SEEK_CUR);
```

- If an **error** occurs (the requested position doesn't exist, for example), `fseek` returns a **nonzero value**.

File positioning

- The file-positioning functions are best used with **binary streams**.
- C doesn't prohibit programs from using them with text streams, but certain restrictions apply.
- For **text streams**, `fseek` can be used only to move to the **beginning** or **end** of a text stream or to return to **a place that was visited previously**.

To remember a position / “bookmark”

- The `ftell` function returns the **current file position** as a **long integer**.
- The value returned by `ftell` may be saved and later supplied to a call of `fseek`:

```
long file_pos;  
...  
file_pos = ftell(fp);  
    /* saves current position */  
...  
fseek(fp, file_pos, SEEK_SET);  
    /* returns to old position */
```

rewind function

- `rewind` function sets the file position at the **beginning**.
- The call `rewind(fp)` is nearly equivalent to `fseek(fp, 0L, SEEK_SET)`.



Example: Copying a file

- The program makes a copy of a file.
- The **names** of the original and new files will be specified on *the command line*.
- An example that uses `fcopy` to copy the file `f1.c` to `f2.c`:

```
fcopy f1.c f2.c
```

- `fcopy` will issue an **error message** if there **aren't exactly two file names** on the command line or if either file **can't be opened**.

Obtaining file names from the command line

- To access **command-line arguments** by defining `main` as a function with two parameters:

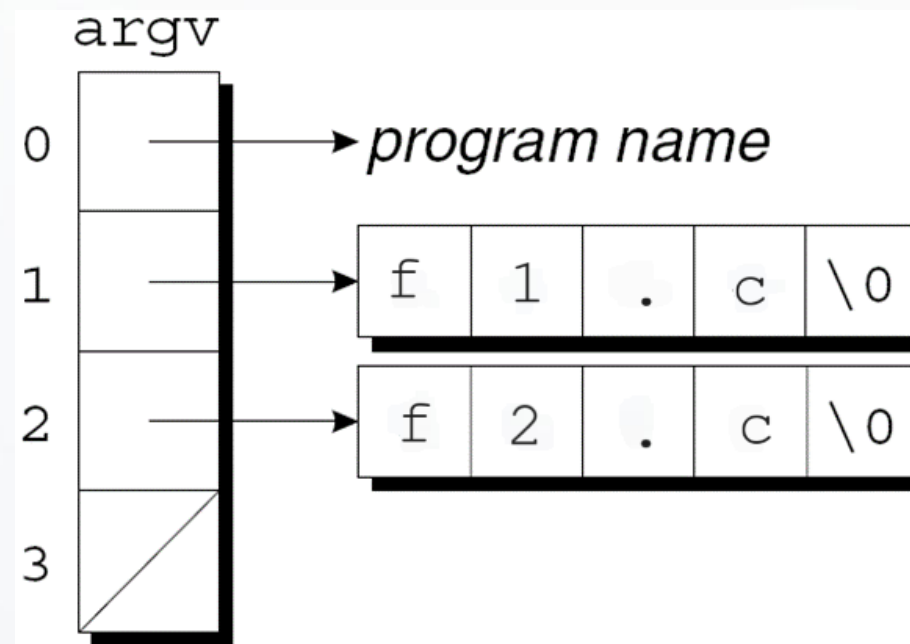
```
int main(int argc, char *argv[])  
{  
    ...  
}
```

- `argc` is the **number of command-line arguments** (including the program name)
- `argv` is an **array of pointers** to the **argument strings**.

Obtaining file names from the command line

- `argv[0]` points to the **program name**
`argv[1]` through `argv[argc-1]` point to the **remaining arguments**
`argv[argc]` is a **null pointer**

```
fcopy f1.c f2.c
```



```
#include <stdio.h>
#include <stdlib.h>
#define BUFSIZE 5000000
char buff[BUFSIZE];
int main(int argc, char *argv[])
{
    FILE *source_fp, *dest_fp;
    if (argc != 3) {
        fprintf(stderr, "usage: fcopy source dest\n");
        exit(EXIT_FAILURE);
    }
    if ((source_fp = fopen(argv[1], "rb")) == NULL) {
        fprintf(stderr, "Can't open %s\n", argv[1]);
        exit(EXIT_FAILURE);
    }
    if ((dest_fp = fopen(argv[2], "wb")) == NULL) {
        fprintf(stderr, "Can't open %s\n", argv[2]);
        fclose(source_fp);
        exit(EXIT_FAILURE);
    }
}
```

```
fseek(source_fp, 0L, SEEK_END);
long int file_length = ftell(source_fp);
if(file_length>BUFSIZE) {
    fprintf(stderr, "File too large to copy!\n");
    fclose(source_fp);
    fclose(dest_fp);
    exit(EXIT_FAILURE);
};

rewind(source_fp);
fread(buff, sizeof(char), file_length, source_fp);
fwrite(buff, sizeof(char), file_length, dest_fp);

fclose(source_fp);
fclose(dest_fp);
return 0;
}
```


Formatted I/O

- `printf` and **related functions** convert data from **binary** form to **text** form during **output**. variable number of arguments

```
int fprintf(FILE *stream,  
            const char *format, ...);
```

return the **number of characters written** (**negative for errors**)

- `scanf` and **related functions** convert data from **text** form to **binary** form during **input**.

```
int fscanf(FILE *stream,  
            const char *format, ...);
```

returns the **number of input items** successfully assigned

Formatted I/O

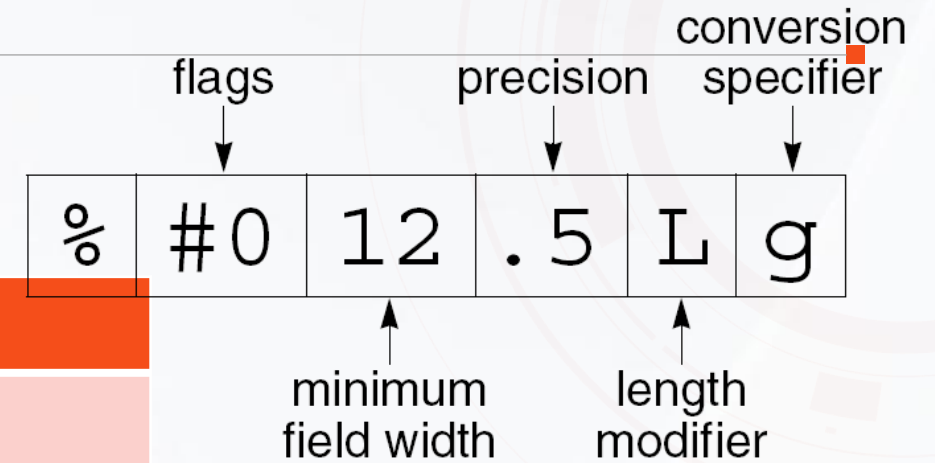
- A call of `printf / scanf` is equivalent to a call of `fprintf / fscanf` with `stdout / stdin` as the first argument.

<code>fprintf(stdout, ...)</code>		<code>printf(...)</code>
<code>fscanf(stdin, ...)</code>		<code>scanf(...)</code>

- Print to `stderr` for error messages to be shown on the screen (even when `stdout` is redirected)

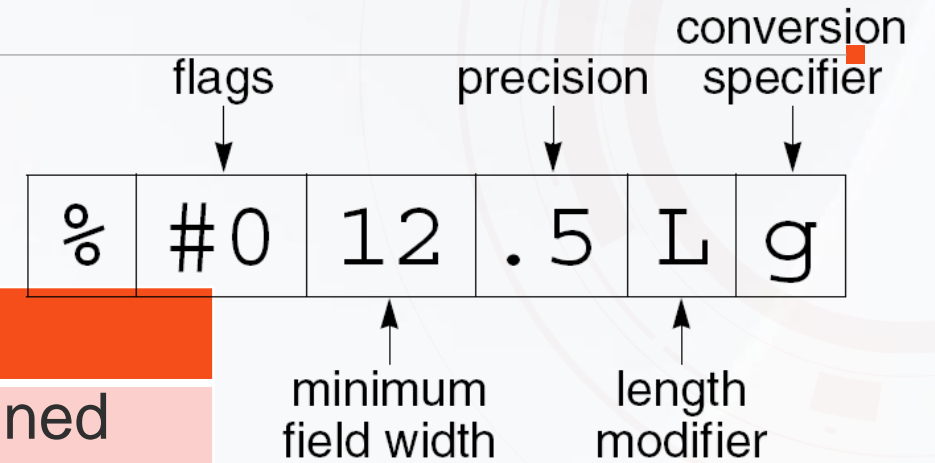
```
fprintf(stderr, "Error: data file can't be  
opened.\n");
```

Conversion specifiers



specifier	Meaning
d, i	an <code>int</code> value in decimal form.
o, u, x, X	an unsigned <code>int</code> value in base 8 (<code>o</code>), base 10 (<code>u</code>), or base 16 (<code>x</code> , <code>X</code>) form. <code>x/X</code> displays digits A–F in lower/upper case.
f	a <code>double</code> value in decimal form, putting the decimal point in the correct position. 6 digits after decimal point by default.
e, E	a <code>double</code> value to scientific notation. 6 digits after decimal point by default.
g, G	<code>g/G</code> converts a <code>double</code> value to either <code>f/F</code> form or <code>e/E</code> form.

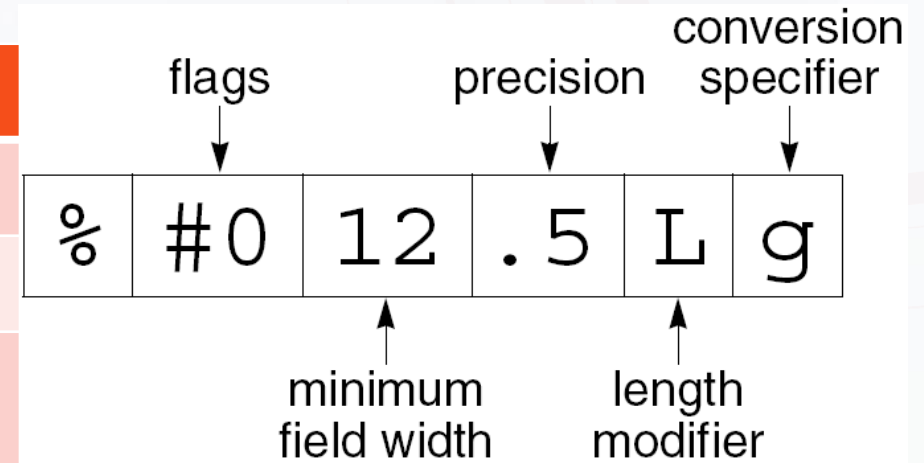
Conversion specifiers



specifier	Meaning
<code>c</code>	Displays an <code>int</code> value as an unsigned character
<code>s</code>	Writes the characters pointed to by the argument. Stops writing when the number of bytes specified by the precision (if present) is reached or a null character is encountered.
<code>p</code>	Print the address in hexadecimal format
<code>n</code>	Stores (no output) the number of characters written so far by this call
<code>%</code>	Writes the character <code>%</code> .

...printf Format Conversion Specifications

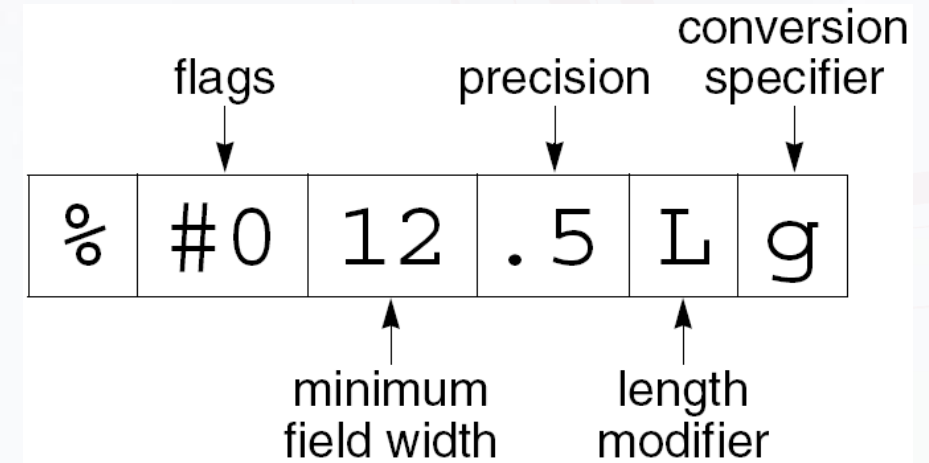
Flags	Meaning
-	Left-justify
+	Signed values always showing +/- sign
space	None negative numbers preceded by as space
#	Octal number begin with 0, hexadecimal numbers begin with 0x or 0X, Floating-point numbers always show decimal points. Trailing zeros aren't removed from numbers printed with the g or G conversions
0	Add leading zeros up to the field width



More than one flag permitted

...printf Format Conversion Specifications

- **Minimum field width** (optional). A shorter item will be padded.
 - By default, **spaces** are added to the **left** of the item.
- A longer item will be displayed completely.
- The field width is either an integer or the character *****.
 - If ***** is present, the field width is obtained from the next argument.



Examples of `...printf` with `%d` conversion

Specification	Result for 123	Result for -123
<code>%8d</code>	<code>.....123</code>	<code>.....-123</code>
<code>%-8d</code>	<code>123.....</code>	<code>-123.....</code>
<code> %+8d</code>	<code>.....+123</code>	<code>.....-123</code>
<code>% 8d</code>	<code>.....123</code>	<code>.....-123</code>
<code>%- 8d</code>	<code>•123.....</code>	<code>-123.....</code>
<code>%-+8d</code>	<code>+123.....</code>	<code>-123.....</code>
<code>%08d</code>	<code>00000123</code>	<code>-0000123</code>
<code> %+08d</code>	<code>+0000123</code>	<code>-0000123</code>
<code>% 08d</code>	<code>•0000123</code>	<code>-0000123</code>

- indicates a space

Effect of the # flag

Specification	Result for 123	Result for 123.0
%8o	•••••173	
%#8o	•••••0173	
%8x	••••••7b	
%#8x	•••••0x7b	
%8X	••••••7B	
%#8X	•••••0X7B	
%8g		•••••123
%#8g		•123.000
%8G		•••••123
%#8G		•123.000

• indicates a space

Examples of the %g conversion

Number	Result of applying % .4g
123456.	1.235e+05
12345.6	1.235e+04
1234.56	1235
123.456	123.5
12.3456	12.35
1.23456	1.235
.123456	0.1235
.0123456	0.01235
.00123456	0.001235
.000123456	0.0001235
.0000123456	1.235e-05
.00000123456	1.235e-06

...printf Format Conversion Specifications

- **Precision** (optional). Meaning depends on the conversion:

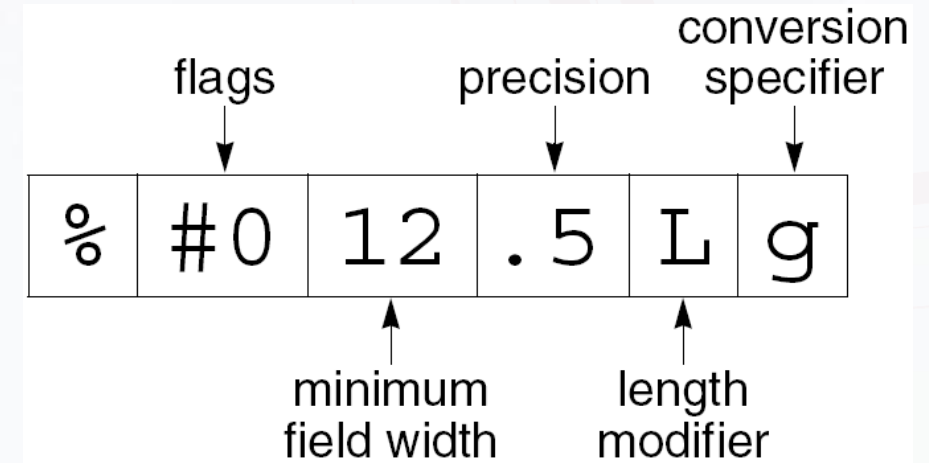
d, i, o, u, x, X: **minimum number of digits**
(leading zeros added if the number has fewer digits)

a, A, e, E, f, F: **number of digits after the decimal point**

g, G: **number of significant digits**

s: maximum **number of bytes**

- The precision is a period (.) followed by an integer or the character * .
 - If * is present, the precision is obtained from the next argument.

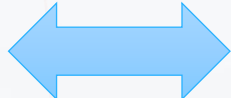


..printf Format Conversion Specifications

Length modifier	Conversion Specifiers	Meaning	%	#0	12	.5	L	g
hh	d, i, o, u, x, X	signed char, unsigned char			↑ minimum field width			↑ length modifier
h	d, i, o, u, x, X	short int, unsigned short int						
l	d, i, o, u, x, X	long int, unsigned long int						
	c	wint_t						
	s	wchar_t *						
ll	d, i, o, u, x, X	long long int, unsigned long long int						
L	a, A, e, E, f, F, g, G	long double						

The ...scanf Functions

- `fscanf` reads data items from an **input stream**, using a **format string** to indicate the input format.
- After the format string, **any number of pointers** follow as additional arguments—each pointing to an input item.
- Input items are converted (according to the **format string**) and stored in these **objects**.

`scanf (...)`  `fscanf(stdin, ...)`

Character I/O

- Library functions read and write single characters (bytes).
- Work with both `text` and `binary` streams.
- The functions treat characters as values of type `int`, not `char`.
- One reason is that the input functions indicate an end-of-file (or error) condition by returning `EOF`, which is a `negative integer constant`.

Output functions

- `fputc` and `putc` write a character to a stream:

```
fputc(ch, fp); /* writes ch to fp */
```

```
putc(ch, fp); Usually a macro
```

- `putchar` writes one character to `stdout`:

```
putchar(ch);  fputc(ch, stdout);
```

Usually a macro

- All three functions set the error indicator for the stream and return `EOF`.
- Otherwise, they return the character that was written.

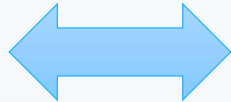
Other input functions

- `fgetc` and `getc` read a **character** from a **stream**:

```
ch = fgetc(fp);
```

```
ch = getc(fp); Usually a macro
```

- All three functions treat the character as an **unsigned char** value (which is then converted to **int** type before it's returned).
- Store the return value in an **int** variable, not a **char**!
- As a result, they never return a **negative** value other than **EOF**.

`getchar()`  `fgetc(stdin)`
Usually a macro